



JAVA Concurrency (SLR 201)

Patrick Bellot & Ada Diaconescu

Télécom ParisTech





Concurrency

- **Concurrency occurs when several tasks are simultaneously executed.**
- **These tasks may interact with each other.**
- **How does a single processor execute several tasks simultaneously ?**
- **Nearly all software requires concurrency.
Example: applications with GUI.**



Concurrent execution

- A sequential program is a series of processor instructions to be executed one after the other.
- A concurrent program consists of several sequential programs, called **processes**, to be executed in parallel.
- On a single processor machine, concurrency is achieved by randomly interleaving the execution of instructions of the processes.
- Processes may affect each other's behavior through shared data and resources, communication, and **synchronization**.



Two types of processes : threads

- Two processes may execute on the same computer sharing memory.
- They are called **threads** or **lightweight processes**.
- Threads can communicate using the shared memory.
- Synchronization problems: access to data in memory, access to resources, etc.



Two types of processes: processes.

- Two processes may execute without sharing memory. They are called **processes** or simply **programs** !
- They may execute on the same computer or on different computers.
- They communicate using **Inter Process Communication (IPC)** resources such as sockets, pipes, etc.
- In JAVA, they can communicate using JAVA/RMI and Sockets.



Main challenges of concurrent programming

- **Synchronizing access to resources.** For instance, two threads must not write in the same file at the same time.
- **Avoiding deadlocks.** If a first thread is waiting for a second thread to do some job and the second thread is waiting for the first one to do some other job, then a deadlock occurs.



Main advantages of concurrency

- **Increased application throughput.** Parallel execution of concurrent processes allows the number of tasks completed in certain time period to increase.
- **High responsiveness for input/output.** I/O intensive applications mostly wait for input or output operations to complete.
- **More appropriate program structure.** Some problems and problem domains are well-suited to representation as concurrent tasks or processes.



Concurrency in JAVA

- Let us study how JAVA implements threads.
- When running a JAVA program, a **thread** is initially created for executing the static main of the program:

```
static void main(String[] args)
```

- Then, the programmer is free to create other threads.



Threads in JAVA

- A **thread** in JAVA is an object of a class inheriting from the **Thread** class.
- The task to be executed is defined by the method `void run()`
- **Practically:**
 - The programmer must create the thread object.
 - Then, it calls its `start()` method to run the thread.
 - When the `run()` method terminates, the thread terminates too.



Example

```
public class MyThread extends Thread
{
    private String threadName ;

    public MyThread(String threadName)
    {
        this.threadName = threadName ;
    }

    public void run()
    {
        for (int i = 0 ; i < 100 ; i++) {
            System.out.println(threadName + ": " + i) ;
            try { Thread.sleep(10) ; } catch(Exception e) {}
        }
    }
}
```



Example (continued)

- To launch such threads (from the `main` method for instance) :

```
MyThread myThreadA = new MyThread("Thread A") ;  
MyThread myThreadB = new MyThread("Thread B") ;
```

```
myThreadB.start() ;  
myThreadA.start() ;
```

```
// Waiting for the two threads to die  
try { myThreadA.join() ; } catch(Exception e) {}  
try { myThreadB.join() ; } catch(Exception e) {}
```



An alternative

- An alternative is to create a class implementing the **Runnable** interface, that means having a **void run()** method.

```
■ public class Task implements Runnable
{
    public void run() { ... .. }
}
```

- Then, to run the task in a thread:

```
Thread thread = new Thread(new Task()) ;
thread.start() ;
```



Sleeping

- The static method `void sleep(...)` of class `Thread` causes the thread executing it to suspend execution for a specified period.
- This is an efficient means of making processor time available to the other threads of an application or to other applications that might be running on the same computer.
- Sleep times are not guaranteed to be precise, because they are limited by the facilities provided by the underlying OS.



Sleeping

■ Syntax:

```
try {  
    Thread.sleep(150) ; // milliseconds  
} catch (Exception e) {}
```

■ Why a try-catch block ?

■ Several `sleep(...)` methods.



Waiting for a thread to terminate

- The method `void join(...)` allows to wait for a thread to terminate:

```
MyThread myThread = new MyThread("...") ;  
myThread.start() ;  
try {  
    myThread.join() ;  
} catch (Exception e) {}
```

- Why a try-catch block ?
- Several `join(...)` methods.



Why try-catch block ?

- A thread can be interrupted using the `interrupt()` method.
- If the thread is currently in a waiting state, for instance sleeping or waiting for another thread or waiting for an I/O, then the waiting function raises an exception.



Pause

- A thread may decide to pause herself using the method `yield()`.
- `yield()` causes the currently executing thread object to temporarily pause and allow other threads to execute.



Need for synchronization

- Interferences between threads occur when two tasks running in different threads but acting on the same data, *interleave*.
- This means that the two tasks consist of multiple steps, and the sequences of steps overlap.
- Interferences on data can be avoided by synchronizing access to this data.



Example

- `int n = 0 ;`
- **Thread A: `n++` ; means:**
 - Get the value of `n` from memory.
 - Increments this value.
 - Stores the new value in memory
- **Thread B: `n--` ; means:**
 - Gets the value of `n` from memory.
 - Decrements this value.
 - Stores the new value in memory
- **Even very simple operations can interleave.**



Synchronized methods

- The methods of an object can be declared as **synchronized**.
- When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object suspend execution until the first thread is done with the object.
- Constructors cannot be synchronized.
- Call to synchronization methods have a cost.



Synchronized methods

```
public class SynchronizedCounter
{
    private int counter ;

    public synchronized void increment() { counter++ ; }

    public synchronized void decrement() { counter-- ; }

    public synchronized int value() { return counter ; }
}
```

With synchronized declarations, method executions cannot interleave.



Synchronized methods

- **JAVA allows to synchronize a method or several statements on a given object.**
- **When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object suspend execution until the first thread is done with the object.**



Synchronized methods

■ Syntax:

```
synchronized (object) {  
    ...  
    instructions ;  
    ...  
}
```



Synchronized methods

- A thread cannot acquire a lock owned by another thread.
- But a thread *can* acquire a lock that it already owns.
- Allowing a thread to acquire the same lock more than once enables **reentrant synchronization**.



Synchronization with wait and notify

- A thread waiting for a condition **MUST NOT** loop:

```
while (! Condition) {} ;
```

- Even with a timer:

```
while (!condition) {  
    try {  
        Thread.sleep(100) ;  
    } catch (Exception e) {}  
}
```

Synchronization with wait and notify

- In JAVA, all objects have a wait-set that contains the set of threads that have executed the `wait()` method of this object.
- `wait()` must be executed in a synchronized block. However, the synchronized lock on the object is released during the wait.
- These threads are sleeping threads waiting to be awoken by a call to the `notify()` or `notifyAll()` method of the object:
 - `notify()` wakes up one thread.
 - `notifyAll()` wakes up all threads.



Example

```
public class CommandsBuffer
{
    // An array to store the commands to be executed
    private String[] commands = new String[1024] ;

    // Index where to store the next arriving command
    // Condition: (nextStoreIdx + 1) % 1024 != lastTakeIdx
    private int nextStoreIdx = 0 ;

    // Index where to take the next command to execute
    // Condition: nextTakeIdx != nextStoreIdx
    private int nextTakeIdx = 0 ;

    ...
}
```



Example continued

```
public synchronized String popCommand()
{
    try {
        while (true) {

            if (nextTakeIdx == nextStoreIdx) {
                wait() ;
            } else {
                String cmd = commands[nextTakeIdx] ;
                nextTakeIdx = (nextTakeIdx + 1) % 1024 ;
                notifyAll() ;
                return cmd ;
            }

        }
    } catch (Exception e) { e.printStackTrace() ;}
    return null ;
}
```



Example continued

```
public synchronized void pushCommand(String cmd)
{
    try {
        while (true) {
            int futureStoreIdx = (nextStoreIdx + 1) % 1024 ;
            if (nextTakeIdx == futureStoreIdx) {
                wait() ;
            } else {
                commands[futureStoreIdx] = cmd ;
                nextStoreIdx = futureStoreIdx ;
                notifyAll() ;
                return ;
            }
        }
    } catch (Exception e) { e.printStackTrace() ;}
}
```



Liveness of threads (SUN Java tutorial)

- **Deadlock** describes a situation where two or more threads are blocked forever, waiting for each other.
- **Starvation** describes a situation where a thread is unable to gain regular access to shared resources and is unable to make progress.
- A thread often acts in response to the action of another thread. If the other thread's action is also a response to the action of another thread, then **livelock** may result.



Conclusions

- **Programming concurrency is tricky.**
- **Be careful !**